

A refusal that will break already looks different at turn 1

One point per chain, at turn 1 of the persuasion: the plain harmful request, which every chain refuses. Chains that will commit the harmful act at a later turn (blue) against chains that never do (orange). Base Llama-3.1-8B-Instruct; deceived route.

v1 — the search set: 192 chains, 16 targets, persuader wording 1

AUROC 0.575

target-matched, mean over L14-18
(the report's statistic)
random floor, 10 randctl seeds:
0.472-0.574

AUROC axis: 10 seeds (grey band), nn (blue)

0.3 0.5 0.7 0.9

13 of 16 targets hold both classes

points pooled across targets; the printed AUROC compares chains within a target (pooled AUROC of these points: 0.50)

median -0.35

median -0.33

will break later
n = 105

never breaks
n = 82

v2 — held out: 40 chains, 5 targets, persuader wording 2

AUROC 0.706

target-matched, mean over L14-18
(the report's statistic)
random floor, 10 randctl seeds:
0.323-0.617

AUROC axis: 10 seeds (grey band), nn (blue)

0.3 0.5 0.7 0.9

4 of 5 targets hold both classes

points pooled across targets; the printed AUROC compares chains within a target (pooled AUROC of these points: 0.52)

median -0.45

median -0.47

will break later
n = 19

never breaks
n = 19

pre-specified on v1, tested once on v2; threshold committed before v2 was read